



PREPARED FOR BELLE REALTY OF LAFAYETTE, LLC · JULY 2026

Orange Ocean Atlas Feature & Decision Brief

Live production system — otb-command.vercel.app · 13 sheets · 269 automated tests · On The Boulevard Shopping Center, Lafayette, LA

Your asset, instrumented: every legal fact, dollar, document, camera, and workflow of On The Boulevard on **one governed surface** — operated by Orange Ocean, with your oversight enforced at the database layer, not by promises.

What the owner personally holds

Standing, read-only truth

Your sign-in sees the same live system the operator runs — dashboard, rent roll, financials — read-only by database policy. Nothing is prepared specially for you, because nothing needs to be: the working system *is* the report.

Why it matters: Owner visibility that doesn't depend on asking.

The monthly Intelligence Brief

On the first of each month: occupancy and scheduled-rent KPIs with real month-over-month deltas, vacancies, the 12-month expiration table — every figure computed, no AI in the loop, archived permanently

in the app.

Why it matters: Fiduciary-grade reporting nobody had to remember to write.

The Owner Safe and the append-only record

Your documents sit in a vault where every view, upload, and delete is logged. From August 1, every dollar of rent lives in an append-only ledger — voids are entries, history is permanent.

Why it matters: If the asset is ever refinanced, audited, or sold, the record is already in evidence-grade shape.

The doctrine behind every decision

Every feature below is a deliberate decision. Four rules govern all of them:

Audit-grade truth.

One authoritative source per fact, with authority rankings, validation rules, and a logged exceptions register. The software stores what was signed and flags what disagrees — it never silently "fixes" data.

Fail closed, never open.

When a safeguard can't verify — an AI number, a rate limit, a signing token, a role — the system declines rather than guesses.

Append-only memory.

Money, compliance, and access events are never edited or deleted; corrections are new entries. The history *is* the audit.

Humans sign; software prepares.

Every outbound consequence — a proposal, a late fee, a signature request — is drafted by the system, released by a person.

The numbers at a glance

MEASURE	VALUE
Asset	62,883 SF GLA · 27 demised units · 2 buildings · 4.84 ac
Occupancy	25 of 27 units (≈95%) · zero holdovers as of July 2026
Income tracked	≈ \$1.08M/yr effective gross, unit-by-unit, PSF-decomposed
Surface	13 sheets · 5 roles · 4 counterparty portals (owner / vendor / tenant / e-sign)
Verification	269 automated tests before every release · production bundle verified per deploy

Instrumentation	17 cameras · 34 parking stalls classified continuously
AI	3 desk personas + 2 phone lines — all grounded, all numerically guardrailed
Build economics	~6-week AI-paired build, single operator · \$0 licensed capture software

1 A governed source of truth

The SOT pack — authority-ranked data governance

Every tenant, rent, and date traces to a vendored source pack with explicit authority ranking (signed rent roll > workbook > any imported system), blocking validation rules, and a known-exceptions log. When a signed figure disagreed with a formula by \$4.84/month, we stored the stated rent and pinned the exception with a test — we did not "fix" it.

Why it matters: Investors underwrite the rent roll. Every figure here answers "says who?" — and the codebase enforces it. Competitors import spreadsheets and hope.

Legal facts as code

The parking variance (324 provided / 344 required), recorded easements, the liquor line, and the anchor's HVAC covenant live as structured data that the dashboard, site plan, and AI leasing screens all derive from — one source, no scattered copies.

Why it matters: The constraints that can kill a deal sit in front of every decision automatically. No competitor models plat-level legal geometry at all.

Confidential seed split

The public app bundle is a skeleton; rents, tenant contacts, and entities hydrate only after role-gated login. Verified in production: zero rent values or tenant emails in public code.

Why it matters: The difference between "password protected" and **protected** — your rent roll isn't downloadable by a visitor.

2 The property twin — plan, 3D, satellite, reality

Plat-exact interactive site plan

The recorded plat as living software: true proportions, easement and liquor lines, 314 drawn stalls, six analytical color lenses, camera and occupancy overlays, operator-placed asset pins.

Why it matters: Owners and buyers see the asset the way a lender's counsel sees it — and the same geometry drives leasing rules, so picture and policy can't drift.

Four registered spatial lenses

Isometric CAD view · orbitable 3D with photogrammetry mesh · georeferenced satellite on a frozen basemap · photoreal drone capture — one shared selection model, built on open tools at \$0 in licenses.

Why it matters: A leasing prospect or buyer tours the asset from a link; alignment is computational and maintainable, not an art project.

Measured parking occupancy — the differentiator

17 cameras feed a hands-free loop: frames every 5 minutes → AI classifies 34 mapped stalls → append-only samples → live paint on the plan + a 7-day dashboard trend, with a heartbeat alert if the pipeline goes quiet.

Why it matters: Parking is this asset's binding legal constraint. Measured utilization is underwriting-grade evidence — and data no competitor in our lane can offer. Everyone else types a number in a box or models it from phone panels.

3 Money & lease economics

Rent roll with one-page owner print

11 columns, PSF decomposition on demand; every bare monthly figure is total rent by rule. The print sheet measures itself to guarantee one page, with unmissable expiry flags (▲ ≤6 months, △ 6–12).

Why it matters: The document an owner forwards to a lender is unambiguous, current, and formatted like an instrument.

Append-only ledger, suggest-only late fees

Charges, payments, voids — all append-only (a void is its own entry). FIFO aging, idempotent monthly charges, late-fee suggestions that post only on operator confirmation. Charges were deliberately gated to August 1 so the system would never fabricate a receivable.

Why it matters: A ledger that reads like a bank statement — reconstructible for any audit, dispute, or sale process.

Deterministic deal engines

Retain-vs-replace and blend-and-extend NER with cap-rate asset-value impact · CAM gross-up · Louisiana eviction sequencing with statutory traps · capital-reserve gating · insurance-claim timelines · an occupancy-cost-ratio tenant-health tripwire. Pure, versioned, unit-tested.

Why it matters: Analyst-grade deal math in seconds, always the same answer for the same question, traceable to formula.

Electronic rent + ACH rail (groundwork live)

Rent is all-electronic by policy; the Stripe webhook and idempotent payment posting are deployed, dormant until armed. Unmatched payments open a manager thread instead of vanishing.

Why it matters: Payments land in the ledger untouched by human keying.

4 AI that may talk but may not guess

The numeric guardrail — crown jewel

Three grounded AI personas answer questions. The moment arithmetic is involved, replies are buffered and every dollar, percentage, and decimal must trace to a deterministic engine output within 0.5% — or the reply is discarded and replaced with the engines' own summary.

Why it matters: The trust story of the product: the AI's numbers **literally cannot be hallucinated**. No AI-native rival makes this posture claim.

Proactive detection, idempotently

A daily scan detects holdovers, 180-day renewal windows, occupancy dips — and opens exactly one AI thread per condition, ever. Proof: it surfaced the 115/117 renewal window early; the tenant renewed on strict terms.

Why it matters: Critical dates stop depending on memory, and automation that cannot duplicate itself is automation an operator can trust.

Monthly Owner Intelligence Brief — deliberately AI-free

A branded monthly report: occupancy and rent KPIs with real month-over-month deltas, vacancies, 12-month expirations, NOI only when a real expense worksheet exists. No language model touches it; generated once, archived immutably.

Why it matters: Fiduciary-grade reporting on the first of every month — without an analyst, and without wondering what an AI embellished.

Voice operations — two phone lines

A 24/7 tenant line that triages by the operator's codified SOPs (emergencies dispatch after-hours; agents never discuss eviction timelines) and files real work orders; a leasing line that screens exclusivity conflicts and books tour slots. Every call becomes a transcript. Built and deployed; carrier go-live in progress.

Why it matters: A missed 2 a.m. leak call is a five-figure loss; a missed leasing call is a lost tenant. Nobody else in the competitive field has voice at all.

5 Counterparty rails

Vendor portal + AI-parsed insurance certificates

Private audited folder and self-service portal per vendor; COI expiry badges; drop an ACORD PDF and AI pre-fills carrier, limits, and expiry for one-click save. COI/W-9/license gate at payment, as policy in the system.

Why it matters: COI lapses are uninsured-loss exposure that lands on the owner. The chase is now systematic and audited.

Event-sourced maintenance

Tenants file requests with photos (portal or phone line) → operator assigns from the trade roster → vendor completes in their portal → every status change is an append-only event; unassigned requests age into alerts.

Why it matters: The maintenance file for any dispute, claim, or diligence is already assembled — who, what, when, with pictures.

In-house e-signatures

Token-lifecycle signing (sent → viewed → signed/declined/expired) with a public branded page, E-SIGN/LA-UETA consent, and database-enforced transitions — double-sign fails closed. The operator delivers the link personally; the app sends nothing on its own.

Why it matters: Signature workflow without a per-envelope SaaS tax, with a verifiable lifecycle for every executed document.

Tenant "Your account" + Owner Safe

Tenants see their own balance from the same append-only ledger the owner sees — scoped by database policy to exactly their unit. Owners get a sealed vault where every view, upload, and delete is logged.

Why it matters: "What do I owe?" stops being a phone call; sensitive owner documents get bank-style access accounting.

6 Security posture & discipline

Defense at the database, verified by audit

Five roles enforced by row-level security that fails closed — a vendor cannot read another vendor's folder even with a modified client. Rate limits fail closed. Secrets live only in the hosting environment with rehearsed, scripted rotation. A dedicated multi-agent security audit ran in July 2026; every code-implementable finding is fixed and deployed.

Why it matters: 269 tests, money-math regression pins, and a deploy rule that verifies the live bundle — the difference between a demo and an operating system for other people's money.

What we deliberately did *not* build

NON-BUILD	REASONING
Our own general ledger	Accounting stays in QuickBooks (sync planned) — we refuse to create an unauditable second GL.
Auto-sent anything	Notices, proposals, signature links: AI drafts, the operator sends. Every outbound consequence is human-released.

Legacy-system import	Evaluated and rejected — its dates were provably unreliable against signed documents. Governed SOT replaced it.
A second codebase	Two sibling builds were harvested for their best engines and archived. One authoritative system.

Competitive position

Verified market scan, July 2026: Yardi/MRI (\$30K–\$150K+/yr, institutional); mid-market generalists with shallow retail depth; and Pickspace, the AI-native direct rival — which has **no spatial layer, no measured occupancy, no voice, no numeric guardrail, and no governed onboarding**. Those are exactly our wedges:

1. **Measured reality** — camera-derived occupancy + a registered spatial twin, vs. KPI tiles and modeled panels.
2. **Numbers that can't hallucinate** — deterministic engines + the fail-closed guardrail.
3. **Governed onboarding** — the SOT authority/validation/exceptions pack as the activation funnel.
4. **Voice-first operations** — phone lines that run the owner's actual SOPs.

Target: owner-operators of 1–20 retail centers — too small for Yardi/Cove, underserved by generalists. Demand already validated: outside owners have asked to get on the software before it has a name.

Where this goes

Phase A (nearly complete): COI AI-parse ✓ · maintenance ✓ · e-sign ✓ · voice built (go-live in progress) · ACH armed for the August 1 ledger go-live. **Phase B:** multi-property schema + free design-partner pilots on validated demand. **Phase C/D:** QuickBooks sync · white-label brand kit · public leasing funnel. OTB remains tenant #1 and the permanent reference dataset — every feature ships proven on a real asset with real money first.

Prepared by **Orange Ocean, LLC** — manager of Belle Realty of Lafayette, LLC · adam@orangeocean.com · orangeocean.ai
Figures as of July 2026. Sources: signed rent roll (docs/sot-2026-07), recorded plat, live production system. This brief describes software built and operated by Orange Ocean; it is not an offer of securities.